

# Signals and signal processing

Suntae Hwang

Kookmin University

# Introduction

---

- ⦿ Signals are software interrupts
- ⦿ Signals provide a way of handling *asynchronous* events
- ⦿ Every signal has a name
  - Begin with the three characters SIG
  - These name are all defined by positive integer constants ( the signal number) in the header <signal.H>
- ⦿ Version 7 had 15 different signals
  - **Unreliable signal model**-get lost and hard to turn off.
- ⦿ SVR4 and 4.3+BSD both have 31 different signals
  - **Reliable signals added.**

# Signal concepts

---

- Numerous conditions can generate a signal
  - The terminal-generated signals occur when user press certain terminal key such as DELETE
  - Hardware exceptions generate signals
    - divide by 0, invalid memory reference and the like
  - The kill(2) function allows a process to send any signal to another process or process group
    - need to be owner of the target process or we have to be a superuser
  - The kill(1) command to send signal to other processes
    - this program is just an interface to the kill function
  - Software conditions can generate signals
    - SIGALRM, SIGPIPE (Broken pipe), SIGURG (Out-of-band data)

# Dispositions of signals

---

## *Disposition or action:*

*Process has to tell the kernel "if and when this signal occurs, do the following."*

### ⦿ Ignore the signal

- This works for most signals, but SIGKILL and SIGSTOP can never be ignored.

### ⦿ Catch the signal

- To do this we tell the kernel to call a function of ours whenever the signal occurs

### ⦿ Let the default action apply

- Every signal has a default action which is to terminate the process in most cases

# Unix signals (ANSI, POSIX.1, SVR4, 4.3+BSD)

---

<b>SIGABRT</b>	abnormal termination(abort)
<b>SIGALRM</b>	time out (alarm)
<b>SIGBUS</b>	hardware fault
<b>SIGCHLD</b>	change in status of a child sent
<b>SIGCONT</b>	continue stopped process
<b>SIGEMT</b>	hardware fault
<b>SIGFPE</b>	arithmetic exception
<b>SIGHUP</b>	hangup
<b>SIGILL</b>	illegal hardware instruction
<b>SIGINFO</b>	status request from keyboard
<b>SIGINT</b>	terminal interrupt character
<b>SIGIO</b>	asynchronous I/O
<b>SIGIOT</b>	hardware fault
<b>SIGKILL</b>	termination
<b>SIGPIPE</b>	write to pipe with no readers
<b>SIGPOLL</b>	pollable event (poll)
<b>SIGPROF</b>	profiling time alarm (setitimer)

<b>SIGPWR</b>	power fail / restart
<b>SIGQUIT</b>	terminal quit character
<b>SIGSEGV</b>	invalid memory reference
<b>SIGSTOP</b>	stop
<b>SIGSYS</b>	invalid system call
<b>SIGTERM</b>	termination
<b>SIGTRAP</b>	hardware fault
<b>SIGTSTP</b>	terminal stop character
<b>SIGTTIN</b>	background read from control tty
<b>SIGTTOU</b>	background write to control tty
<b>SIGURG</b>	urgent condition
<b>SIGUSR1</b>	user-defined signal
<b>SIGUSR2</b>	user-defined signal
<b>SIGVTALRM</b>	virtual time alarm (setitimer)
<b>SIGWINCH</b>	terminal window size change
<b>SIGXCPU</b>	CPU limit exceeded
<b>SIGXFSZ</b>	file size limit exceeded

# Signals

---

- ◉ SIGART :
  - generated by calling the `abort` function.
- ◉ SIGALRM :
  - generated when a timer set with the `alarm` expires.
- ◉ SIGCHLD :
  - Whenever a process terminates or stops, the signal is sent to the parent.
- ◉ SIGCONT :
  - This signal(job-control) sent to a stopped process when it is continued.
- ◉ SIGFPE :
  - signals an arithmetic exception, such as divide-by-0, floating point overflow, and so on
- ◉ SIGHUP :
  - generated to the controlling process (session leader) associated with a controlling terminal if a disconnect is detected by the terminal interface
  - generated if the session leader terminates and sent to each process in the foreground process group
  - commonly used to notify daemon process to reread their configuration files (note that a daemon should not have a controlling terminal and normally never receive this signal)

# Signals (cont'd)

---

- ◉ SIGILL :
  - indicates that the process has executed an illegal hardware instruction.
- ◉ SIGINT :
  - generated by the terminal driver when we type the interrupt key and sent to all processes in the foreground process group
- ◉ SIGIO :
  - indicates an asynchronous I/O event
- ◉ SIGKILL :
  - can't be caught or ignored. a sure way to kill any process.
- ◉ SIGPIPE :
  - If we write to a pipeline but the reader has terminated, SIGPIPE is generated
- ◉ SIGPWR :
  - related to power failure. (read the book for the detail)
- ◉ SIGQUIT :
  - generated by the terminal driver when we type terminal quit key and sent to all processes in the foreground process group

# Signals (cont'd)

---

- ◉ SIGSEGV :
  - indicates that their process has made an invalid memory reference
- ◉ SIGSTOP :
  - This signal(job-control) stops a process and can't be caught or ignored
- ◉ SIGSYS :
  - signals an invalid system call
- ◉ SIGTERM :
  - the termination signal sent by the `kill(1)` command by default.
- ◉ SIGTSTP :
  - This is the interactive stop signal generated by the terminal driver when we type the terminal suspend key and sent to all processes in the foreground process group.
- ◉ SIGTTIN :
  - generated by the terminal driver when a process in a background process group tries to read from its controlling terminal
- ◉ SIGTTOU :
  - generated by the terminal driver when a process in a background process group tries to write to its controlling terminal



# Signals (cont'd)

---

## • SIGURG :

- notifies the process that an urgent condition has occurred. Optionally generated when out-of-band data is received on a network connection.

## • SIGUSR1[2] :

- user-defined signals, for use in application programs

## • SIGWINCH :

- generated to the foreground process group when a process changes the window size from its previous value, with the `ioctl` set-window-size command

## • SIGXCPU :

- generated if the process exceeds its soft CPU time limit

## • SIGXFSZ :

- generated if the process exceeds its soft file size limit

# Signal Function

---

```
#include <signal.h>
```

```
void ( *signal( int signo, void (*func) (int))) (int)
```

Returns: previous disposition of signal if OK, SIG\_ERR on error

- ① The simplest interface to the signal features of Unix
    - *signo* : the name of the signal
    - *func* :
      - SIG\_IGN - ignore the signal
      - SIG\_DFL - take its default action
      - The address of a signal handler ( or signal-catching function): a function to be called (catching) when the signal occurs.
    - The signal handler is passed a single integer argument (*the signal number*) and returns nothing.
    - signal() returns the pointer to the previous signal handler
- ```
typedef void Sigfunc(int);  
Sigfunc *signal(int, Sigfunc *);
```

# Signal Function Example

```
static void sig_child(int);
int main(void) {
    pid_t pid; int i;
    signal(SIGCHLD, sig_child);

    pid = fork();
    if (pid == 0) {
        sleep(1);
        exit(0);
    }
    while(1) { i = i; }
}

static void
sig_child(int signo){
    pid_t pid; int status;
    pid = wait(&status);
    printf("child %d finished\n", pid);
}
```

```
$ a.out
child 17145 finished
```

```
static void sig_fpe(int);
int main(void) {
    pid_t pid; int i;
    signal(SIGFPE, sig_fpe);
    i = i/0;
}

static void
sig_fpe(int signo){
    pid_t pid; int status;
    printf("Divide by 0 Error\n");
    /* routine that saves all variables*/
    exit(1);
}
```

```
$ a.out
Floating point exception
```

```
$ a.out
Divide by 0 Error
```

# Signal Function Example

```
#include <signal.h>
static void sig_usr(int); /* one handler for both
signals */
int main(void){
    if (signal(SIGUSR1, sig_usr) == SIG_ERR)
        err_sys("can't catch SIGUSR1");
    if (signal(SIGUSR2, sig_usr) == SIG_ERR)
        err_sys("can't catch SIGUSR2");
    for ( ; ; ) pause();
}
static void
sig_usr(int signo) { /* argument is signal number */
    if (signo == SIGUSR1)
        printf("received SIGUSR1\n");
    else if (signo == SIGUSR2)
        printf("received SIGUSR2\n");
    else err_dump("received signal %d\n", signo);
    return;
}
```

```
$ a.out &
[1] 4720
$ kill -USR1 4720 send it SIGUSR1
received SIGUSR1
$ kill -USR2 4720 send it SIGUSR2
received SIGUSR2
```

```
$kill 4720 send it SIGTERM
[1] + Terminated a.out &
```

# Program Start-up

---

- ◉ When a process is forked, the child inherits the parent's signal dispositions.
- ◉ When a program is executed
  - the disposition of any signals that are being caught to their default action
  - the status of all other signals (ignored or default) is left alone
- ◉ An interactive shell (w/o job control)
  - sets the disposition of the interrupt and quit signals in the background process to be ignored
  - Many interactive programs catches the signals only when not in the background (the signal is not ignored) by doing the following:

```
int sig_int(), sig_quit()  
if (signal(SIGINT, SIG_IGN) != SIG_IGN) signal(SIGINT, sig_int);  
if (signal(SIGQUIT, SIG_IGN) != SIG_IGN) signal(SIGQUIT, sig_quit);
```

# Interrupted System Calls (1/2)

---

## • Slow system calls : that can block forever

- reads from/writes to files that can block the caller forever (pipes, terminal, network)
- open files that block until some condition occurs (opening terminal devices that waits until a modem answers the phone)
- pause() and wait()
- certain ioctl() operations and some IPC functions

## • A slow system call is interrupted by a signal

- returns an error and errno was set to EINTR
- need to handle the error explicitly

Again:

```
if ((n = read(fd, buff, BUFSIZE)) < 0) {  
    if (errno == EINTR) go to Again; /* interrupted system call */  
}
```

# Interrupted System Calls (2/2)

---

- Automatic restarting of certain interrupted system calls (4.2BSD)
  - ioctl, read, readv, write, writev, wait and waitpid (wait, waitpid are always interrupted when a signal is caught)
  - 4.3BSD allow to disable this feature on a per-signal basis
  - Without the automatic restart feature, we need to test every read/write for the interrupted error return and reissue the read or write.
- Fast system calls completes before the signal was delivered

# Reentrant Functions

---

- ① **POSIX.1 specifies the functions that are guaranteed to be reentrant**
- ① **Calling a none-reentrant function from a signal handler may produce unpredictable results**
  - While the main program calls malloc() and interrupted, the signal handler also calls malloc(), then what could happen?
- ① **One errno variable per process even with reentrant guaranteed functions - save the errno and restore it later.**



# Reentrant functions that may be called from a signal handler

|                          |                        |                          |                          |
|--------------------------|------------------------|--------------------------|--------------------------|
| <code>_exit</code>       | <code>fork</code>      | <code>pipe</code>        | <code>stat</code>        |
| <code>abort*</code>      | <code>fstat</code>     | <code>read</code>        | <code>sysconf</code>     |
| <code>access</code>      | <code>getegid</code>   | <code>rename</code>      | <code>tcdrain</code>     |
| <code>alarm</code>       | <code>geteuid</code>   | <code>rmdir</code>       | <code>tcflow</code>      |
| <code>cfgetispeed</code> | <code>getgid</code>    | <code>setgid</code>      | <code>tcflush</code>     |
| <code>cfgetospeed</code> | <code>getgroups</code> | <code>setpgid</code>     | <code>tcgetattr</code>   |
| <code>cfsetispeed</code> | <code>getpgrp</code>   | <code>setsid</code>      | <code>tcgetpgrp</code>   |
| <code>cfsetospeed</code> | <code>getpid</code>    | <code>setuid</code>      | <code>tcsendbreak</code> |
| <code>chdir</code>       | <code>getppid</code>   | <code>sigaction</code>   | <code>tcsetattr</code>   |
| <code>chmod</code>       | <code>getuid</code>    | <code>sigaddset</code>   | <code>tcsetpgrp</code>   |
| <code>chown</code>       | <code>kill</code>      | <code>sigdelset</code>   | <code>time</code>        |
| <code>close</code>       | <code>link</code>      | <code>sigemptyset</code> | <code>times</code>       |
| <code>creat</code>       | <code>longjmp*</code>  | <code>sigfillset</code>  | <code>umask</code>       |
| <code>dup</code>         | <code>lseek</code>     | <code>sigismember</code> | <code>uname</code>       |
| <code>dup2</code>        | <code>mkdir</code>     | <code>signal*</code>     | <code>unlink</code>      |
| <code>execle</code>      | <code>mkfifo</code>    | <code>sigpending</code>  | <code>utime</code>       |
| <code>execve</code>      | <code>open</code>      | <code>sigprocmask</code> | <code>wait</code>        |
| <code>exit*</code>       | <code>pathconf</code>  | <code>sigsuspend</code>  | <code>waitpid</code>     |
| <code>fcntl</code>       | <code>pause</code>     | <code>sleep</code>       | <code>write</code>       |

# Reentrant Functions (cont'd)

```
err_sys(char *s) { fprintf(stderr,"%s",s); exit(1);}
static void my_alarm(int);

int main(void) {
    struct passwd *ptr;
    signal(SIGALRM, my_alarm); alarm(1);
    for ( ; ; ) {
        if ( (ptr = getpwnam("sthwang")) == NULL) err_sys("getpwnam error");
        if (strcmp(ptr->pw_name, "sthwang") != 0)
            printf("return value corrupted!, pw_name = %s\n", ptr->pw_name);
    }
}

static void my_alarm(int signo) {
    struct passwd *rootptr;

    printf("in signal handler\n");
    if ( (rootptr = getpwnam("root")) == NULL)
        err_sys("getpwnam(root) error");
    alarm(1);
    return;
}
```

**\$ a.out**

in signal handler  
Segmentation fault

**\$ a.out**

in signal handler  
in signal handler  
Segmentation fault

**\$ a.out**

in signal handler  
getpwnam(root) error

# Kill and Raise function (1/2)

---

```
#include <sys/types.h>
#include <signal.h>
int kill(pid_t pid, int signo);
int raise(int signo);
Both return: 0 if OK, 1 on error
```

- ① The kill function sends a signal to a process or a group of process
  - pid > 0 signal to the process whose process ID is pid
  - pid == 0 signal to the processes whose process group ID equals that of sender
  - pid < 0 signal to the processes whose process group ID equals abs. of pid
  - pid == -1 POSIX.1 leaves this condition unspecified (used as a broadcast signal in SVR4, 4.3+BSD)
- ② The raise function allows a process to send a signal to itself

# Kill and Raise function (2/2)

---

- **A process needs permission to send a signal to some other process**
  - The superuser can send a signal to any process
  - The real or effective user ID of the sender has to equal the real or effective user ID of the receiver
  - SIGCONT can be sent to any member process of the same session
  - **signo = 0: null signal**
    - normal error checking performed, but no signal is sent
    - used often to determine if a specific process still exists. (If the process doesn't exist, kill returns -1 and errno is set to ESRCH).

# alarm and pause function (1/2)

---

```
#include <unistd.h>
unsigned int alarm (unsigned int seconds) ;
    Returns: 0 or number of seconds until previously set alarm
```

## ● Alarm function

- sets a timer that will expire at a specified time in the future
- When the timer expires, the SIGALRM signal generated
- seconds is the number of clock seconds in the future when the signal should be generate
- default action of the signal is to terminate the process.
- There could be a extra delay between when the signal generated and when the signal handler gets the control
- only one alarm clock per process
  - previously registered alarm clock is replaced by the new value
  - if seconds=0, the previous alarm clock is cancelled

# alarm and pause function (2/2)

---

```
#include <unistd.h>
int pause (void) ;
```

**Returns:** -1 with errno set to EINTR

## ● Pause function

- suspends the calling process until a signal is caught.
- returns only if a signal handler is executed and that handler returns.
- returns -1 with errno set to EINTR

# Example I (sleep1)

```
static void
sig_alm(int signo)
{
    return; /* nothing to do, just return to wake up the pause */
}

unsigned int
sleep1(unsigned int nsecs) /* sleep the process for nsecs */
{
    if (signal(SIGALRM, sig_alm) == SIG_ERR)
        return(nsecs);
    alarm(nsecs); /* start the timer */
    pause(); /* next caught signal wakes us up */
    return( alarm(0) ); /* turn off timer, return unslept time */
}
```

- If the caller of sleep1() already has an alarm set, the alarm is erased by the first call to alarm.
  - Save remaining alarm time and reset the alarm before the return
- Modify the disposition for SIGALRM
  - Save the disposition and reset before the return
- Race condition: alarm may goes off before the pause(); the caller is suspended forever at pause()=> sigpromask, sigsuspend

## Example II (sleep2)

---

```
static jmp_buf env_alm;
static void
sig_alm(int signo)
{
    longjmp(env_alm, 1);
}
unsigned int
sleep2(unsigned int nsecs)
{
    if (signal(SIGALRM, sig_alm) == SIG_ERR)
        return(nsecs);
    if (setjmp(env_alm) == 0) {
        alarm(nsecs);           /* start the timer */
        pause();                /* next caught signal wakes us up */
    }
    return( alarm(0) );         /* turn off timer, return unslept time */
}
```

- The previous race condition was avoided
- Another problem if SIGALRM interrupts some other signal handler and the longjmp() aborts the other signal handler (see the next example)




# Example III (sleep2 problem)

```
int main(void){
    unsigned int unslept;
    if (signal(SIGINT, sig_int) == SIG_ERR)
        err_sys("signal(SIGINT) error");
    unslept = sleep2(5);
    printf("sleep2 returned: %u\n", unslept);
    exit(0);
}

static void
sig_int(int signo){ /* the for loop executes more than 5 sec */
    int i;
    volatile int j;
    printf("\nsig_int starting\n");
    for (i = 0; i < 2000000; i++) j += i * i;
    printf("sig_int finished\n");
    return;
}
```

longjmp



\$ a.out

^?

sig\_int starting

sleep2 returned: 0

sleep2 starts running

Type our interrupt char

SIGALRM generated while in sig\_int()

longjmp aborted sig\_int

## Example IV (timeout)

```
int main(void){
    int n; char line[MAXLINE];
    if (signal(SIGALRM, sig_alm) == SIG_ERR)
        err_sys("signal(SIGALRM) error");
    alarm(10);
    if ( (n = read(STDIN_FILENO, line, MAXLINE)) < 0)
        err_sys("read error");
    alarm(0);
    write(STDOUT_FILENO, line, n);
    exit(0);
}
static void
sig_alm(int signo){
    return; /* nothing to do, just return to interrupt the read */
}
```

- ⦿ A common use for alarm : timeout function
- ⦿ Race condition: alarm may go off before read()
- ⦿ If the read system call is automatically restarted, timeout does not work.

# Example V (Another timeout)

```
static jmp_buf env_alm;
int
main(void){
    int n;char line[MAXLINE];
    if (signal(SIGALRM, sig_alm) == SIG_ERR)
        err_sys("signal(SIGALRM) error");
    if (setjmp(env_alm) != 0)
        err_quit("read timeout");
    alarm(10);
    if ( (n = read(STDIN_FILENO, line, MAXLINE)) < 0)
        err_sys("read error");
    alarm(0);
    write(STDOUT_FILENO, line, n);
    exit(0);
}
static void
sig_alm(int signo)
{
    longjmp(env_alm, 1);
}
```

- ⦿ No problems with automatic restart
- ⦿ But still has the race condition and the problem with other signal handler interactions

# Abort Function

---

```
#include <stdlib.h>
void abort(void);
```

**This function never returns**

- ⦿ Causes abnormal program termination
- ⦿ This function sends the SIGABRT signal to the process
- ⦿ SIGABRT signal handler to perform any cleanup that it wants to do, before the process terminated
- ⦿ POSIX.1 states that if the process does not terminate itself from this signal handler, when signal handler returns, abort terminates the process.

# Sleep Function

---

```
#include <signal.h>
unsigned int sleep(unsigned int seconds) ;
Returns: 0 or number of unslept seconds
```

- ⦿ This function causes the calling process to be suspended until either
  - The amount of wall clock time specified by second has elapsed
    - The return value is 0
  - A signal is caught by the process and the signal handler returns
    - The return value is the number of unslept seconds
  - The actual return may be at a time later than requested because of other system activity
  - There can be interactions between sleep and alarm if sleep is implemented with the alarm functions (unspecified by POSIX.1)

# Unreliable signals (1/2)

---

- ⦿ Signals were unreliable in earlier version of UNIX (V7)
  - Signals could get lost
    - a signal could occur and the process would never know about it
  - The action for the signal was reset to its default each time the signal occurred
  - A process had a little control over a signal
    - can catch or ignore the signal, but can't block it
    - unable to turn a signal off when it didn't want the signal to occur (all it can do was ignore the signal)
    - can't "prevent the following signals from occurring, but only remember if they do occur".

# Unreliable signals(2/2)

```
int sig_int(); /*my signal handling function*/
...
signal(SIGINT, sig_int); /*establish handler*/
...
sig_int()
{
    signal(SIGINT, sig_int);
    /*reestablish handler for next occurrence*/
    ... /*process the signal*/
}
```

```
/* set the flag to remember that an interrupt occurs
*/
int sig_int_flag;
main() {
    int sig_int();
    ...
    signal(SIGINT, sig_int);
    ...
    while(sig_int_flag==0)
        pause(); /* go to sleep, waiting for signal */
    ...
}
sig_int() {
    signal(SIGINT, sig_int);
    sig_int_flag=1; /*set flag for main loop to
examine*/
}
```

**Q:** What happens if the interrupt signal occurs again before `signal()` in `sig_int()` executes?

**Q:** What happens if the interrupt signal occurs again before `pause()` and after `while` statement?

# Reliable Signal Terminology and Semantics

---

- ⦿ A signal is *generated* for a process (or sent to a process) when the event that causes the signal occurs
  - when signal is generated the kernel usually sets a flag of some form in the process table
- ⦿ A signal is *delivered* to a process when the action for a signal is taken
- ⦿ During the time between the generation of signal and its delivery, the signal is said to be *pending*.
- ⦿ Each process has a *signal mask* that defines the set of signals currently *blocked* from delivery to that process



# Reliable Signal Terminology and Semantics

---

- ⦿ *blocking* the delivery of a signal
  - A blocked signal (with default or catch signal action) remains pending until
    - unblocks the signal
    - changes the action to ignore the signal
- ⦿ – What to do with a blocked signal is determined when the signal is delivered, not when it is generated
  - This allows the process to change the action for the signal before it is delivered
  - Most Unix does not queue blocked signals generated more than once; the Unix kernel just delivers the signal once.

# Signal sets

---

- ⦿ A data type to represent multiple signals (sigset\_t)
- ⦿ Do not assume global/static variable initialization in C for sigset\_t
- ⦿ Functions to manipulate signal sets

```
#include <signal.h>
int sigemptyset (sigset_t *set);
int sigfillset (sigset_t *set);
int sigaddset (sigset_t *set, int signo);
int sigdelset (sigset_t *set, int signo);
all four return: 0 if OK, -1 on error
int sigismember ( const sigset_t *set, int signo);
Returns:1 if true, 0 if false
```

# Signal Functions

---

- ◉ Superset of the functionality of `signal()`
- ◉ `sigprocmask(int how, const sigset_t * setp, sigset_t * osetp);`
  - examine or change signal masks
- ◉ `sigpending(sigset_t* setp);`
  - Return the set of signals that are blocked and pending
- ◉ `sigaction(int signo, const struct sigaction* act, struct sigaction oact);`
  - examine or modify the action associated with a particular signal
- ◉ `sigsuspend(const sigset_t* sigmask);`
  - atomic operation (reset the signal mask and pause )
- ◉ `sigsetjmp(sigjmp_buf env, int savemask)/siglongjmp(sigjmp_buf env, init val)`
  - nonlocal branching from a signal handler

# Sigprocmask function

```
#include <signal.h>
int sigprocmask(int how, const sigset_t *set, sigset_t *oset);
Returns: 0 if OK, -1 on error
```

- ⦿ Examine or change the *signal mask* that is the set of signals currently blocked from delivery to the process
  - First, if *oset* is nonnull pointer, the current signal mask for the process is returned through *oset*
  - Second, if *set* is a nonnull pointer, then the *how* argument indicates how the current signal mask is modified
  - if *set* is NULL, *how* is not significant
  - *how*
- ⦿ SIG\_BLOCK : redefine the new signal mask (*set + oset*)
- ⦿ SIG\_UNBLOCK : unblock the signals in *set*
- ⦿ SIG\_SETMASK : the new signal mask = *set*

if there are any pending, unblocked signals after `sigprocmask` at least one of these signal is delivered to the process before `sigprocmask` returns

# Sigprocmask Function Example

---

```
#include <errno.h>
#include <signal.h>
#include "ourhdr.h"

void
pr_mask(const char *str)
{
    sigset_t sigset;
    int errno_save;
    errno_save = errno; /* we can be called by signal handlers */
    if (sigprocmask(0, NULL, &sigset) < 0)
        err_sys("sigprocmask error");
    printf("%s", str);
    if (sigismember(&sigset, SIGINT)) printf("SIGINT ");
    if (sigismember(&sigset, SIGQUIT)) printf("SIGQUIT ");
    if (sigismember(&sigset, SIGUSR1)) printf("SIGUSR1 ");
    if (sigismember(&sigset, SIGALRM)) printf("SIGALRM ");
    /* remaining signals can go here */
    printf("\n");
    errno = errno_save;
}
```

# Sigpending Function

---

```
#include <signal.h>
int sigpending (sigset_t *set);
Returns: 0 if OK, -1 on error
```

- ⦿ Returns the set of signals that are blocked from delivery and currently pending for the calling process
  - The set of signals is returned through the set argument

# Sigpending Function

```
static void sig_quit(int);
int main(void) {
    sigset_t newmask, oldmask, pendmask;
    if (signal(SIGQUIT, sig_quit) == SIG_ERR)
        err_sys("can't catch SIGQUIT");
    sigemptyset(&newmask); sigaddset(&newmask, SIGQUIT);

    /* block SIGQUIT and save current signal mask */
    if (sigprocmask(SIG_BLOCK, &newmask, &oldmask) < 0)
        err_sys("SIG_BLOCK error");
    sleep(5); /* SIGQUIT here will remain pending */

    if (sigpending(&pendmask) < 0)
        err_sys("sigpending error");
    if (sigismember(&pendmask, SIGQUIT))
        printf("\nSIGQUIT pending\n");

    /* reset signal mask which unblocks SIGQUIT */
    if (sigprocmask(SIG_SETMASK, &oldmask, NULL) < 0)
        err_sys("SIG_SETMASK error");
    printf("SIGQUIT unblocked\n");
    sleep(5); /* SIGQUIT here will terminate with core file */
    exit(0);
}
```

# Sigpending Function

```
static void
sig_quit(int signo)
{
    printf("caught SIGQUIT\n");

    if (signal(SIGQUIT, SIG_DFL) == SIG_ERR)
        err_sys("can't reset SIGQUIT");
    return;
}
```

**\$ a.out**

**^\\^\\^\\^\\**

**SIGQUIT pending**

**caught SIGQUIT**

**SIGQUIT unblocked**

**^\\Quit(coredump)**

generate signal several times(before  
5 seconds are up)

after return from sleep(5)

in signal handler (signal is generated once!)

after return from sigprocmask

generate signal again (message by shell)



# Sigaction Function (1/3)

---

```
#include <signal.h>
int sigaction (int signo, const struct sigaction *act,
struct sigaction *oact);
```

**Returns: 0 if OK, -1 on error**

- ⦿ Examine or modify the action associated with a particular signal -  
supersedes the signal function from earlier UNIX
  - If act pointer is nonnull, we are modifying the action.
  - If oact pointer is nonnull, the system returns the previous action for the signal.
- ⦿ Specify a set of signals added to the signal mask before the  
signal handler is called - *includes the signal being delivered*.
  - This way we are able to *block* certain signals whenever a signal handler invoked
  - When the signal-catching function returns, the signal mask of process is *reset* to its previous value
- ⦿ The action remains installed until we explicitly change it by  
calling sigaction().

# Sigaction Function (2/3)

---

```
struct sigaction {  
    void (*sa_handler)(int);          /* address of signal handler */  
    sigset_t sa_mask;                 /* additional signals to block */  
    int sa_flags;                     /* signal options */  
    void (*sa_sigaction)(int, siginfo_t *, void *);  
};
```

- sa\_handler: points to the signal handler or SIG\_IGN or SIG\_DFL
- sa\_mask: additional signals to block when the signal handler (as opposed to SIG\_IGN or SIG\_DFL) is called
- sa\_flags : specifies various options for the handling of the signal
  - SA\_RESTART - system calls interrupted are automatically restarted
  - SA\_NODEFER - when this signal is caught, the signal is not automatically blocked by the system while the signal handler executes (unreliable signal)
  - SA\_RESETHAND - the disposition for the signal is reset to SIG\_DFL on entry to the signal handler (unreliable signal)
  - SA\_SIGINFO – provides additional information to a signal handler. Final sig\_action field is used

# Sigation Function (3/3)

```
/* An implementation of reliable signal() using sigaction */
/* for unreliable signal(), use SA_RESETHAND and SA_NODEFER */

#include <signal.h>
Sigfunc *
signal(int signo, Sigfunc *func) {
    struct sigaction act, oact;

    act.sa_handler = func;
    sigemptyset(&act.sa_mask);
    act.sa_flags = 0;

    if (signo != SIGALRM) {
        act.sa_flags |= SA_RESTART; /*SVR4, 4.3+BSD*/
    }

    if (sigaction(signo, &act, &oact) < 0)
        return(SIG_ERR);

    return (oact.sa_handler);
}
```

# Sigsetjmp and siglongjmp Functions

---

```
#include <set jmp.h>
int sigsetjmp (sigjmp_buf env, int savemask) ;
    Returns: 0 if called directly, nonzero if returning from a call to siglongjmp
int siglongjmp (sigjmp_buf env, int val) ;
```

- What happens to the signal mask for the process if we longjmp out of the signal handler?
- sigsetjmp and siglongjmp saves and restores the signal mask
  - use these functions for nonlocal branching from a signal handler
- If *savemask* is nonzero then sigsetjmp also saves the current signal mask of the process in *env*
- When siglongjmp is called, if the *env* argument was saved with nonzero *savemask*, then siglongjmp restores the saved signal mask

# Sigsetjmp and siglongjmp Functions

<Example of sigsetjmp and siglongjmp>

```
static sigjmp_buf jmpbuf;
static volatile sig_atomic_t canjump;
int main(void) {
    ...
    if(sigsetjmp(jmpbuf, 1)) exit(0);
    canjump=1;                                /* now sigsetjmp() is OK */
    ...
}
static void sig_usr1(int signo) {
    ...
    if (canjump==0) return;                  /* unexpected signal, */
   /* not jmpbuf ready, ignored */
    ...
    canjump = 0;
    siglongjmp(jmpbuf, 1);
}
```

## ☉ □ Use of canjump variable

- protection against the signal handler being called when the jump buffer isn't initialized by sigsetjmp
- sig\_atomic\_t : variable can be written without being interrupted (ex. No page boundary crossing )

# Sigsuspend Function (1/4)

---

```
#include <signal.h>
int sigsuspend (const sigset_t *sigmask) ;
Returns: -1 with errno set to EINTR
```

- ⦿ Performs resetting the signal mask and put the process to sleep in a single atomic operation
  - Signal mask of the process is set to the value pointed to by sigmask.
  - The process is also suspended until a signal is caught or until a signal occurs that terminates the process
  - If a signal is caught and if the signal handler returns, then `sigsuspend` returns and the signal mask of the process is set to its old value

## Sigsuspend Function (2/4)

```
/* protect critical regions of code from interrupt signals : a wrong way */
sigset_t newmask, oldmask;

sigemptyset(&newmask);
sigaddset(&newmask, SIGINT);
    /* block SIGINT and save current signal mask */
if (sigprocmask(SIG_BLOCK, &newmask, &oldmask) < 0)
    err_sys("SIG_BLOCK error");
    /* critical region of code */
    /* reset signal mask which unblocks SIGINT */
if (sigprocmask(SIG_SETMASK, &oldmask, NULL) < 0)
    err_sys("SIG_SETMASK error");
pause(); /* wait for signal to occur */
    /* and continue processing ... */
```

- ⦿ Problem: any signal between the second sigprocmask() and pause() gets lost.

# Sigsuspend Function (3/4)

```
/* protect critical regions of code from interrupt signals : a right way */
sigset_t newmask, oldmask, zeromask;
if (signal(SIGINT, sig_int) == SIG_ERR) err_sys("signal(SIGINT) error");

sigemptyset(&zeromask);
sigemptyset(&newmask);
sigaddset(&newmask, SIGINT);
/* block SIGINT and save current signal mask */
if (sigprocmask(SIG_BLOCK, &newmask, &oldmask) < 0) err_sys("SIG_BLOCK error");

/* critical region of code */
/* allow all signals and pause */
if (sigsuspend(&zeromask) != -1) err_sys("sigsuspend error");

/* reset signal mask which unblocks SIGINT */
if (sigprocmask(SIG_SETMASK, &oldmask, NULL) < 0) err_sys("SIG_SETMASK error");

/* and continue processing ... */
```

- ⦿ Eliminated the previous problems. (unblock and pause)
- ⦿ Note: need the second sigprocmask() to unblock SIGINT because the return of sigsuspend() set the signal mask to its value before the call.



# Sigsuspend Function (4/4)

```
volatile sig_atomic_t quitflag; /* set nonzero by signal handler */
int main(void) {
    sigset_t newmask, oldmask, zeromask;
    if (signal(SIGINT, sig_int) == SIG_ERR) err_sys("signal(SIGINT) error");
    if (signal(SIGQUIT, sig_int) == SIG_ERR) err_sys("signal(SIGQUIT) error");

    sigemptyset(&zeromask); sigemptyset(&newmask);
    sigaddset(&newmask, SIGQUIT); /* block SIGQUIT and save current signal mask */

    if (sigprocmask(SIG_BLOCK, &newmask, &oldmask) < 0) err_sys("SIG_BLOCK error");

    while (quitflag == 0) sigsuspend(&zeromask);

    /* SIGQUIT has been caught and is now blocked; do whatever */
    quitflag = 0;
    if (sigprocmask(SIG_SETMASK, &oldmask, NULL) < 0) err_sys("SIG_SETMASK error");
    exit(0);
}

void sig_int(int signo) { /* one signal handler for SIGINT and SIGQUIT */
    if (signo == SIGINT) printf("\ninterrupt\n");
    else if (signo == SIGQUIT) quitflag = 1; /* set flag for main loop */
    return;
}
```

- ④ sigsuspend to wait for a global variable to be set.

# Sleep(1/2)

---

```
static void sig_alm(void) {
    return; /* nothing to do, just returning wakes up sigsuspend() */
}
unsigned int sleep(unsigned int nsecs) {
    struct sigaction newact, oldact;
    sigset_t newmask, oldmask, suspmask;
    unsigned int unslept;

    newact.sa_handler = sig_alm;
    sigemptyset(&newact.sa_mask);
    newact.sa_flags = 0;
    sigaction(SIGALRM, &newact, &oldact);
    /* set our handler, save previous information */

    sigemptyset(&newmask);
    sigaddset(&newmask, SIGALRM);
    /* block SIGALRM and save current signal mask */
    sigprocmask(SIG_BLOCK, &newmask, &oldmask);
```

# Sleep(2/2)

---

```
alarm(nsecs);

suspmask = oldmask;
sigdelset(&suspmask, SIGALRM); /* make sure SIGALRM isn't blocked */

sigsuspend(&suspmask); /* wait for any signal to be caught */

        /* some signal has been caught, SIGALRM is now blocked */

unslept = alarm(0);
sigaction(SIGALRM, &oldact, NULL); /* reset previous action */

        /* reset signal mask, which unblocks SIGALRM */
sigprocmask(SIG_SETMASK, &oldmasks, NULL);
return(unslept);
}
```

- Handles signals reliably
- Avoiding the race condition
- Do not handle any interactions with previously set alarms